

Sign Language Detection

Artificial Intelligence, MCSC102
Department of Computer Science,
University of Delhi

Harsh Yadav (20)
Purnima Kumar (41)
Yashi Sharma (58)



Overview

01 Quick Recap

02 Methodology

03 Deep Neural Networks

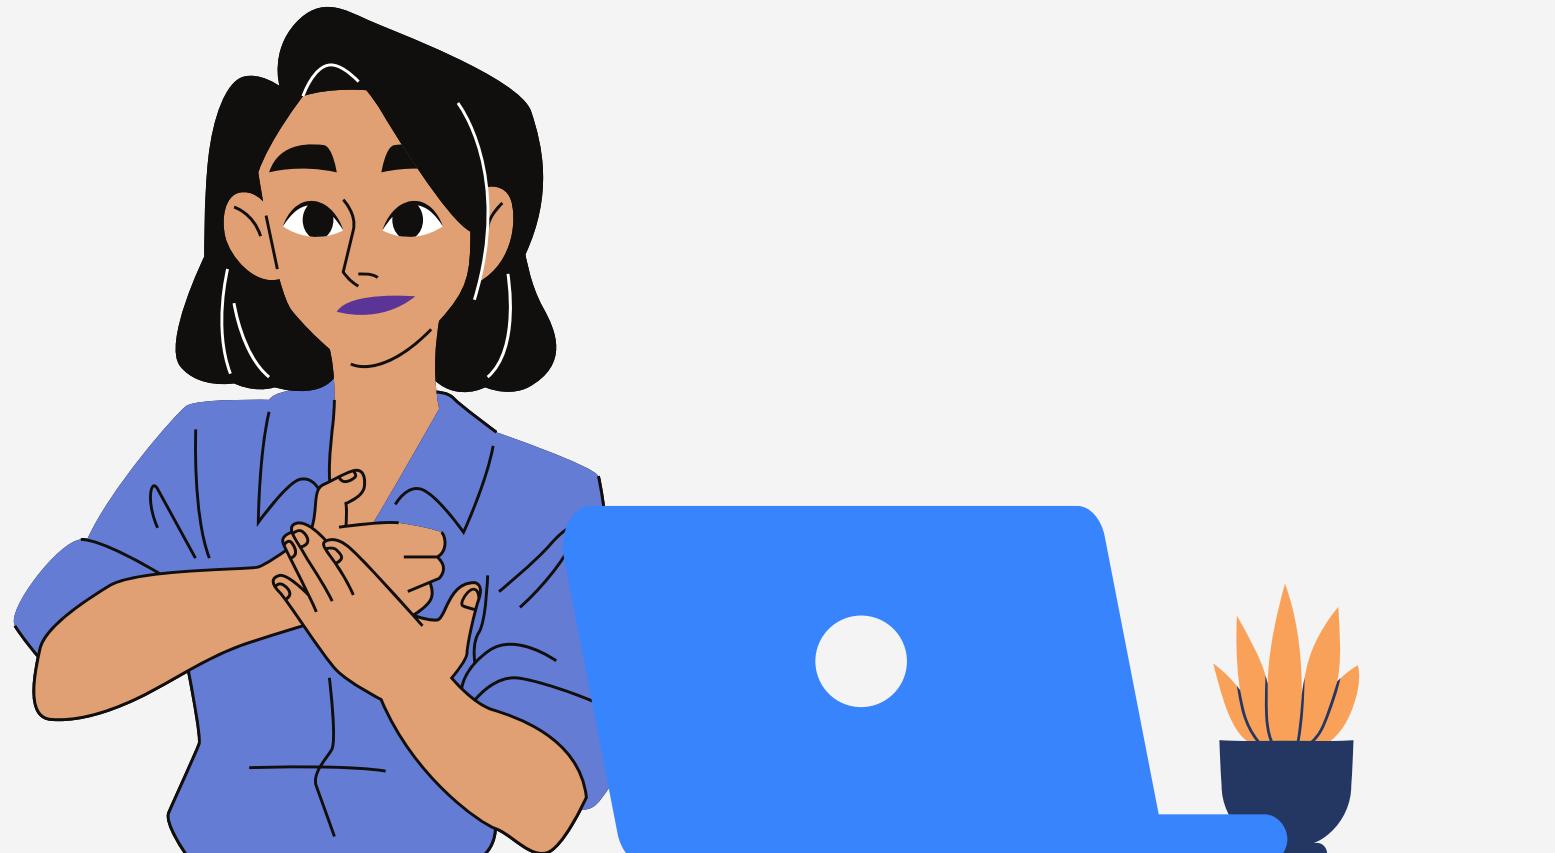
04 Transfer Learning

05 Preliminary Experiment Results

06 References



Quick Recap



- We aim to develop a **computer-vision based AI system** that can accurately interpret and translate **sign language gestures**.
- **Input:** Images of hand gestures representing American Sign Language (ASL).
- **Output:** Accurate textual translations of English alphabet (A-Z) and numbers (0-9).

Methodology

Having conducted a comprehensive literature review, we have identified certain methodologies that have demonstrated notable efficacy in the detection of Sign Language Gestures.

Transfer Learning

What is transfer learning?

It is a process where the weights of an existing pretrained model that was created to perform a similar task are retrained on a new data in order to perform a new, but similar task.



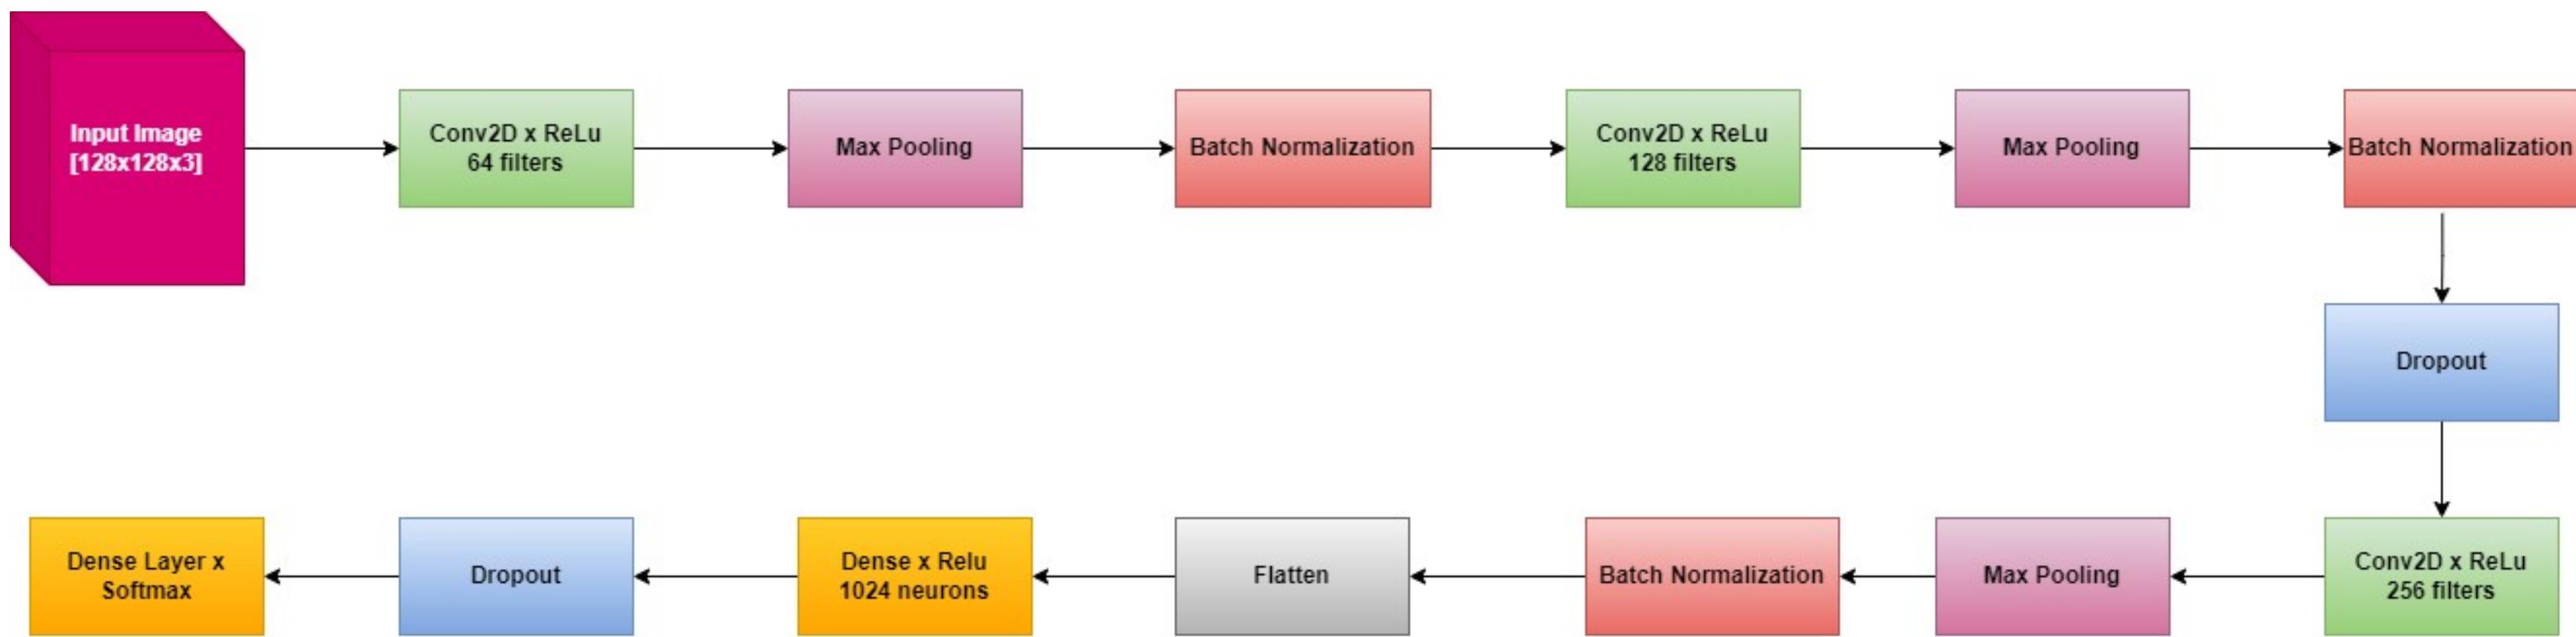
Convolutional Neural Networks

Since, we are dealing with a multi-class image classification problem Convolutional Neural Networks is a go-to architecture.

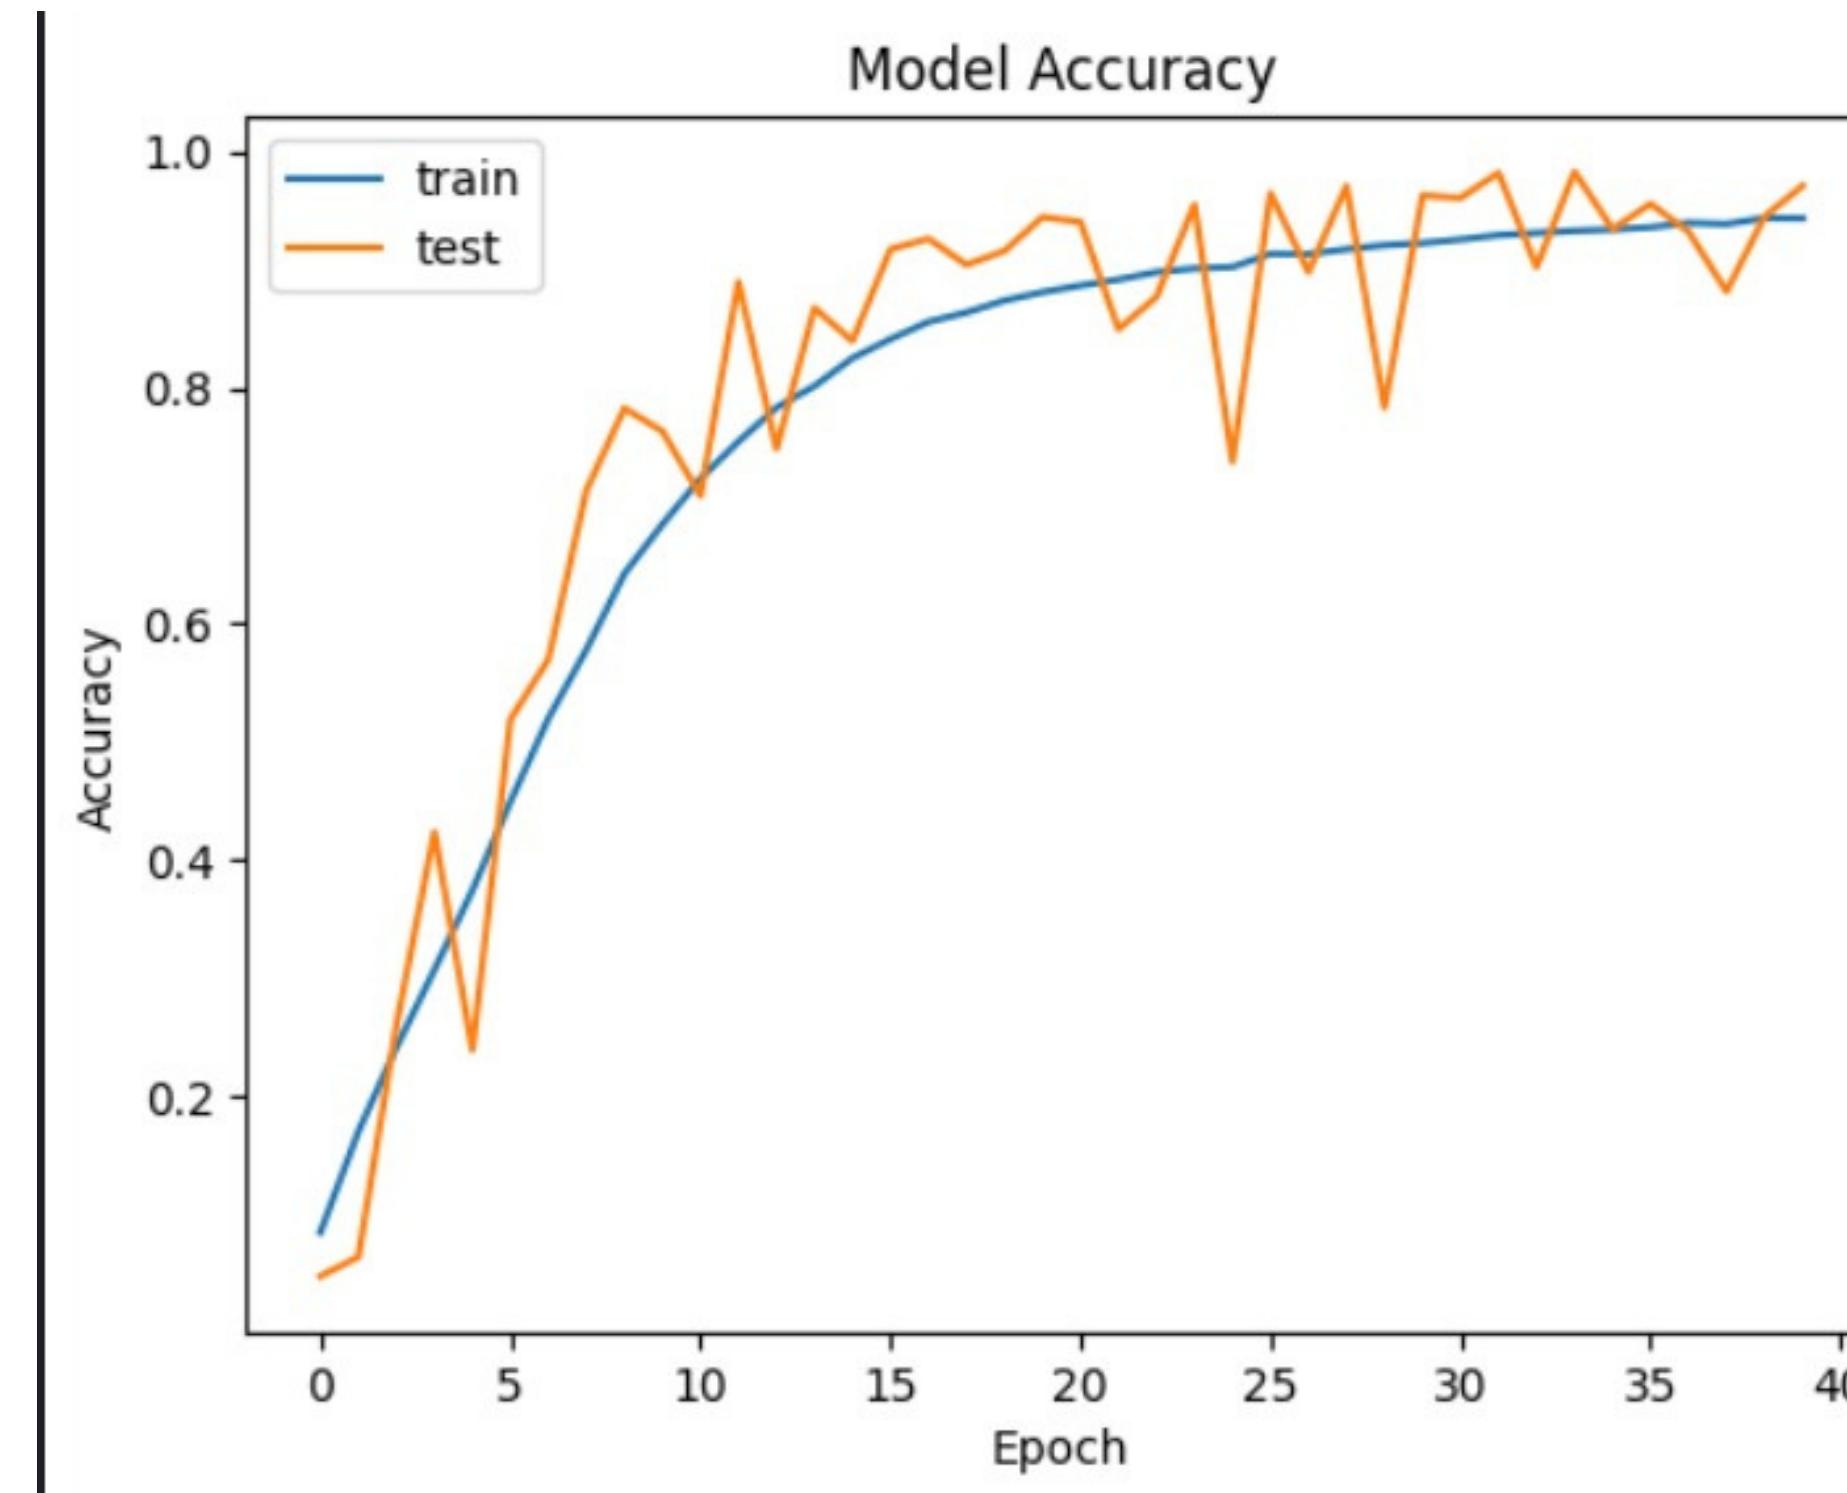
CNN Architecture

- 01 Input**
INPUT holds the raw pixel values of the image, in this case an image of width , height , and with three color R, G, B.
- 02 Convolutional Layers**
The Convolution layers computes the output of the neurons that are connected to local regions in the input.
- 03 Pooling Operations**
The Pooling layer performs a downsampling operation along with the spatial dimensions.
- 04 Activation Functions**
ReLU activation function applies an elementwise activation function, i.e., $f(x) = \max(0, x)$.
- 05 Fully Connected Layers**
In the end, the fully connected layers computes the class scores.
- 06 Output**
We get a single neuron as an output which contains class scores for all the classes.

Implementation Using CNNs

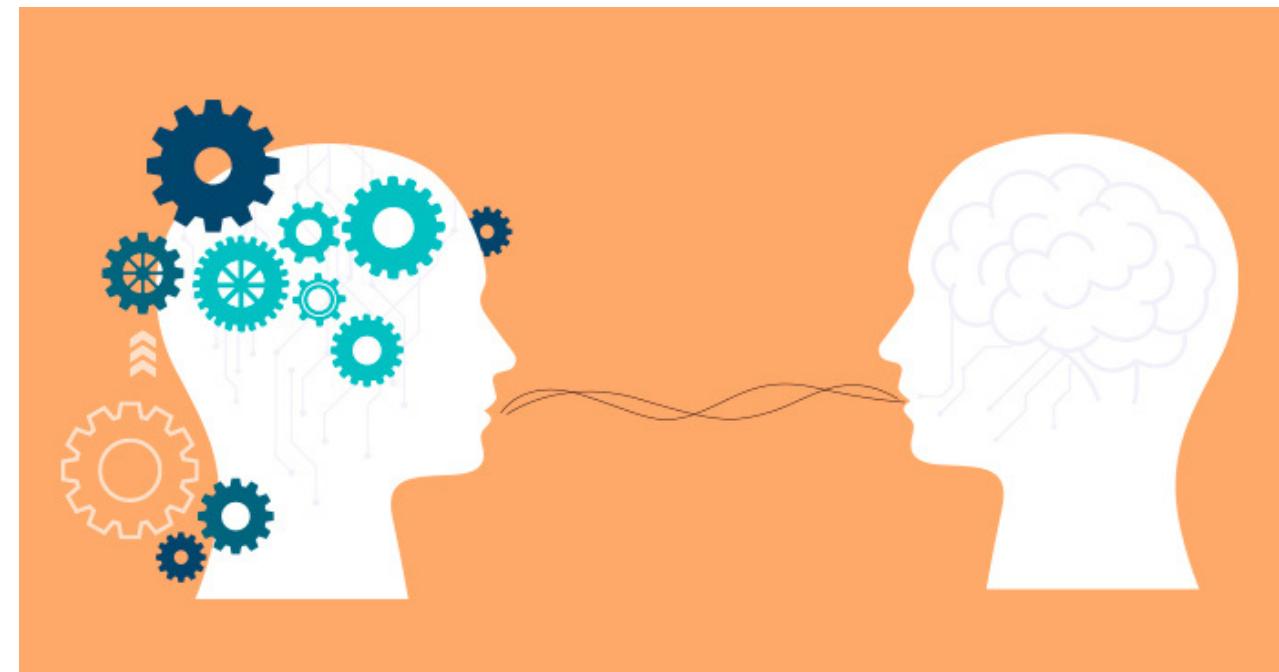


Experimental Results



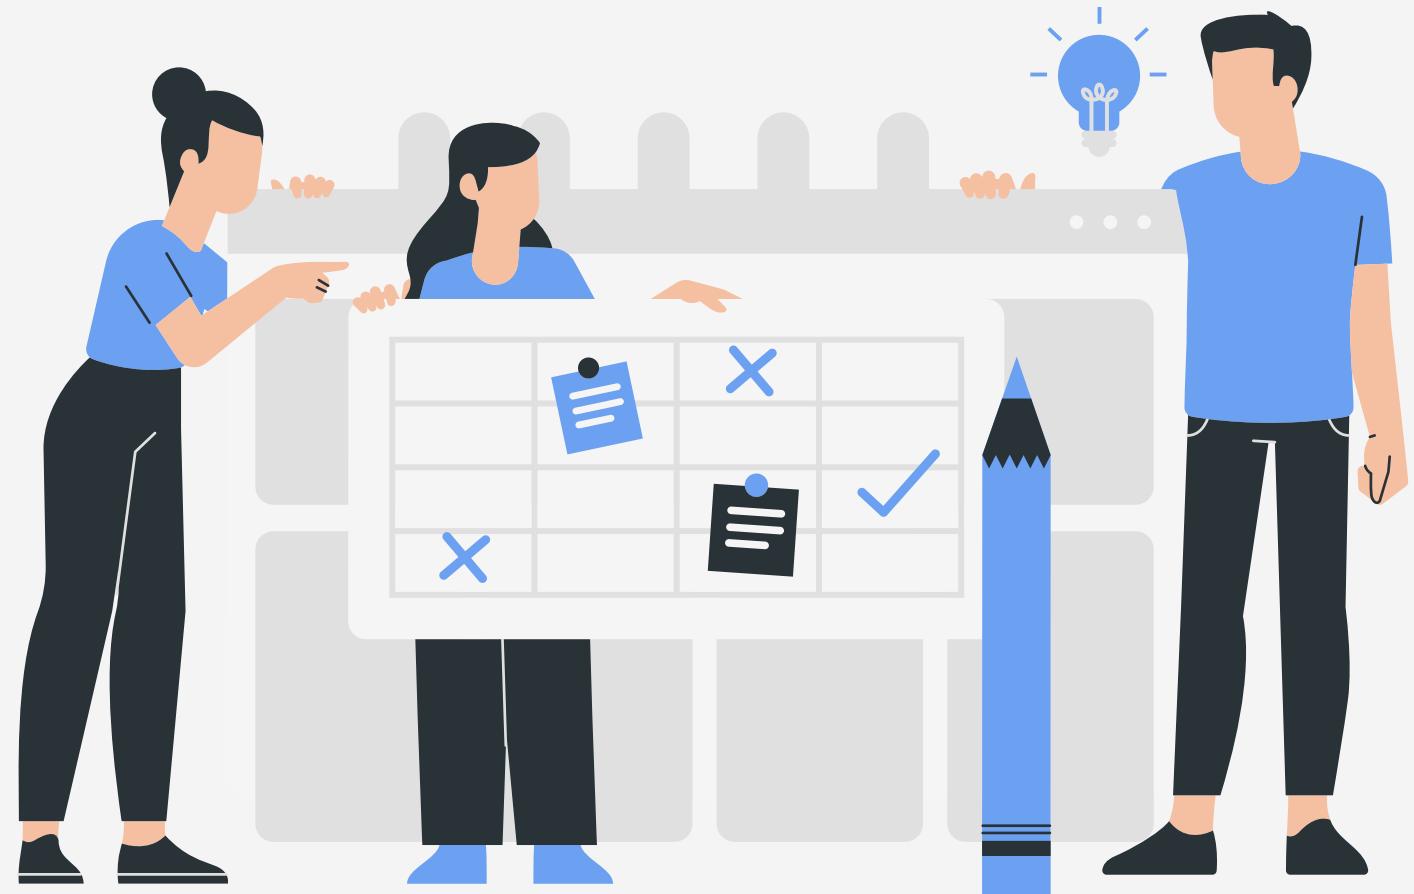
Transfer Learning Approach

- Transfer Learning is a process where the **weights of an existing pretrained model** that was created to perform a similar task are **retrained on a new data** in order to perform a new, but similar task.



- Used Microsoft COCO dataset, which is a large-scale object detection dataset containing 330,000 images consisting of 80 object categories.

Methodology applied by us to detect words.



Captured Images using OpenCV

5 classes (10 images each)
Hello, Yes, No, ThankYou, ILoveYou



Hello

YES



I Love You



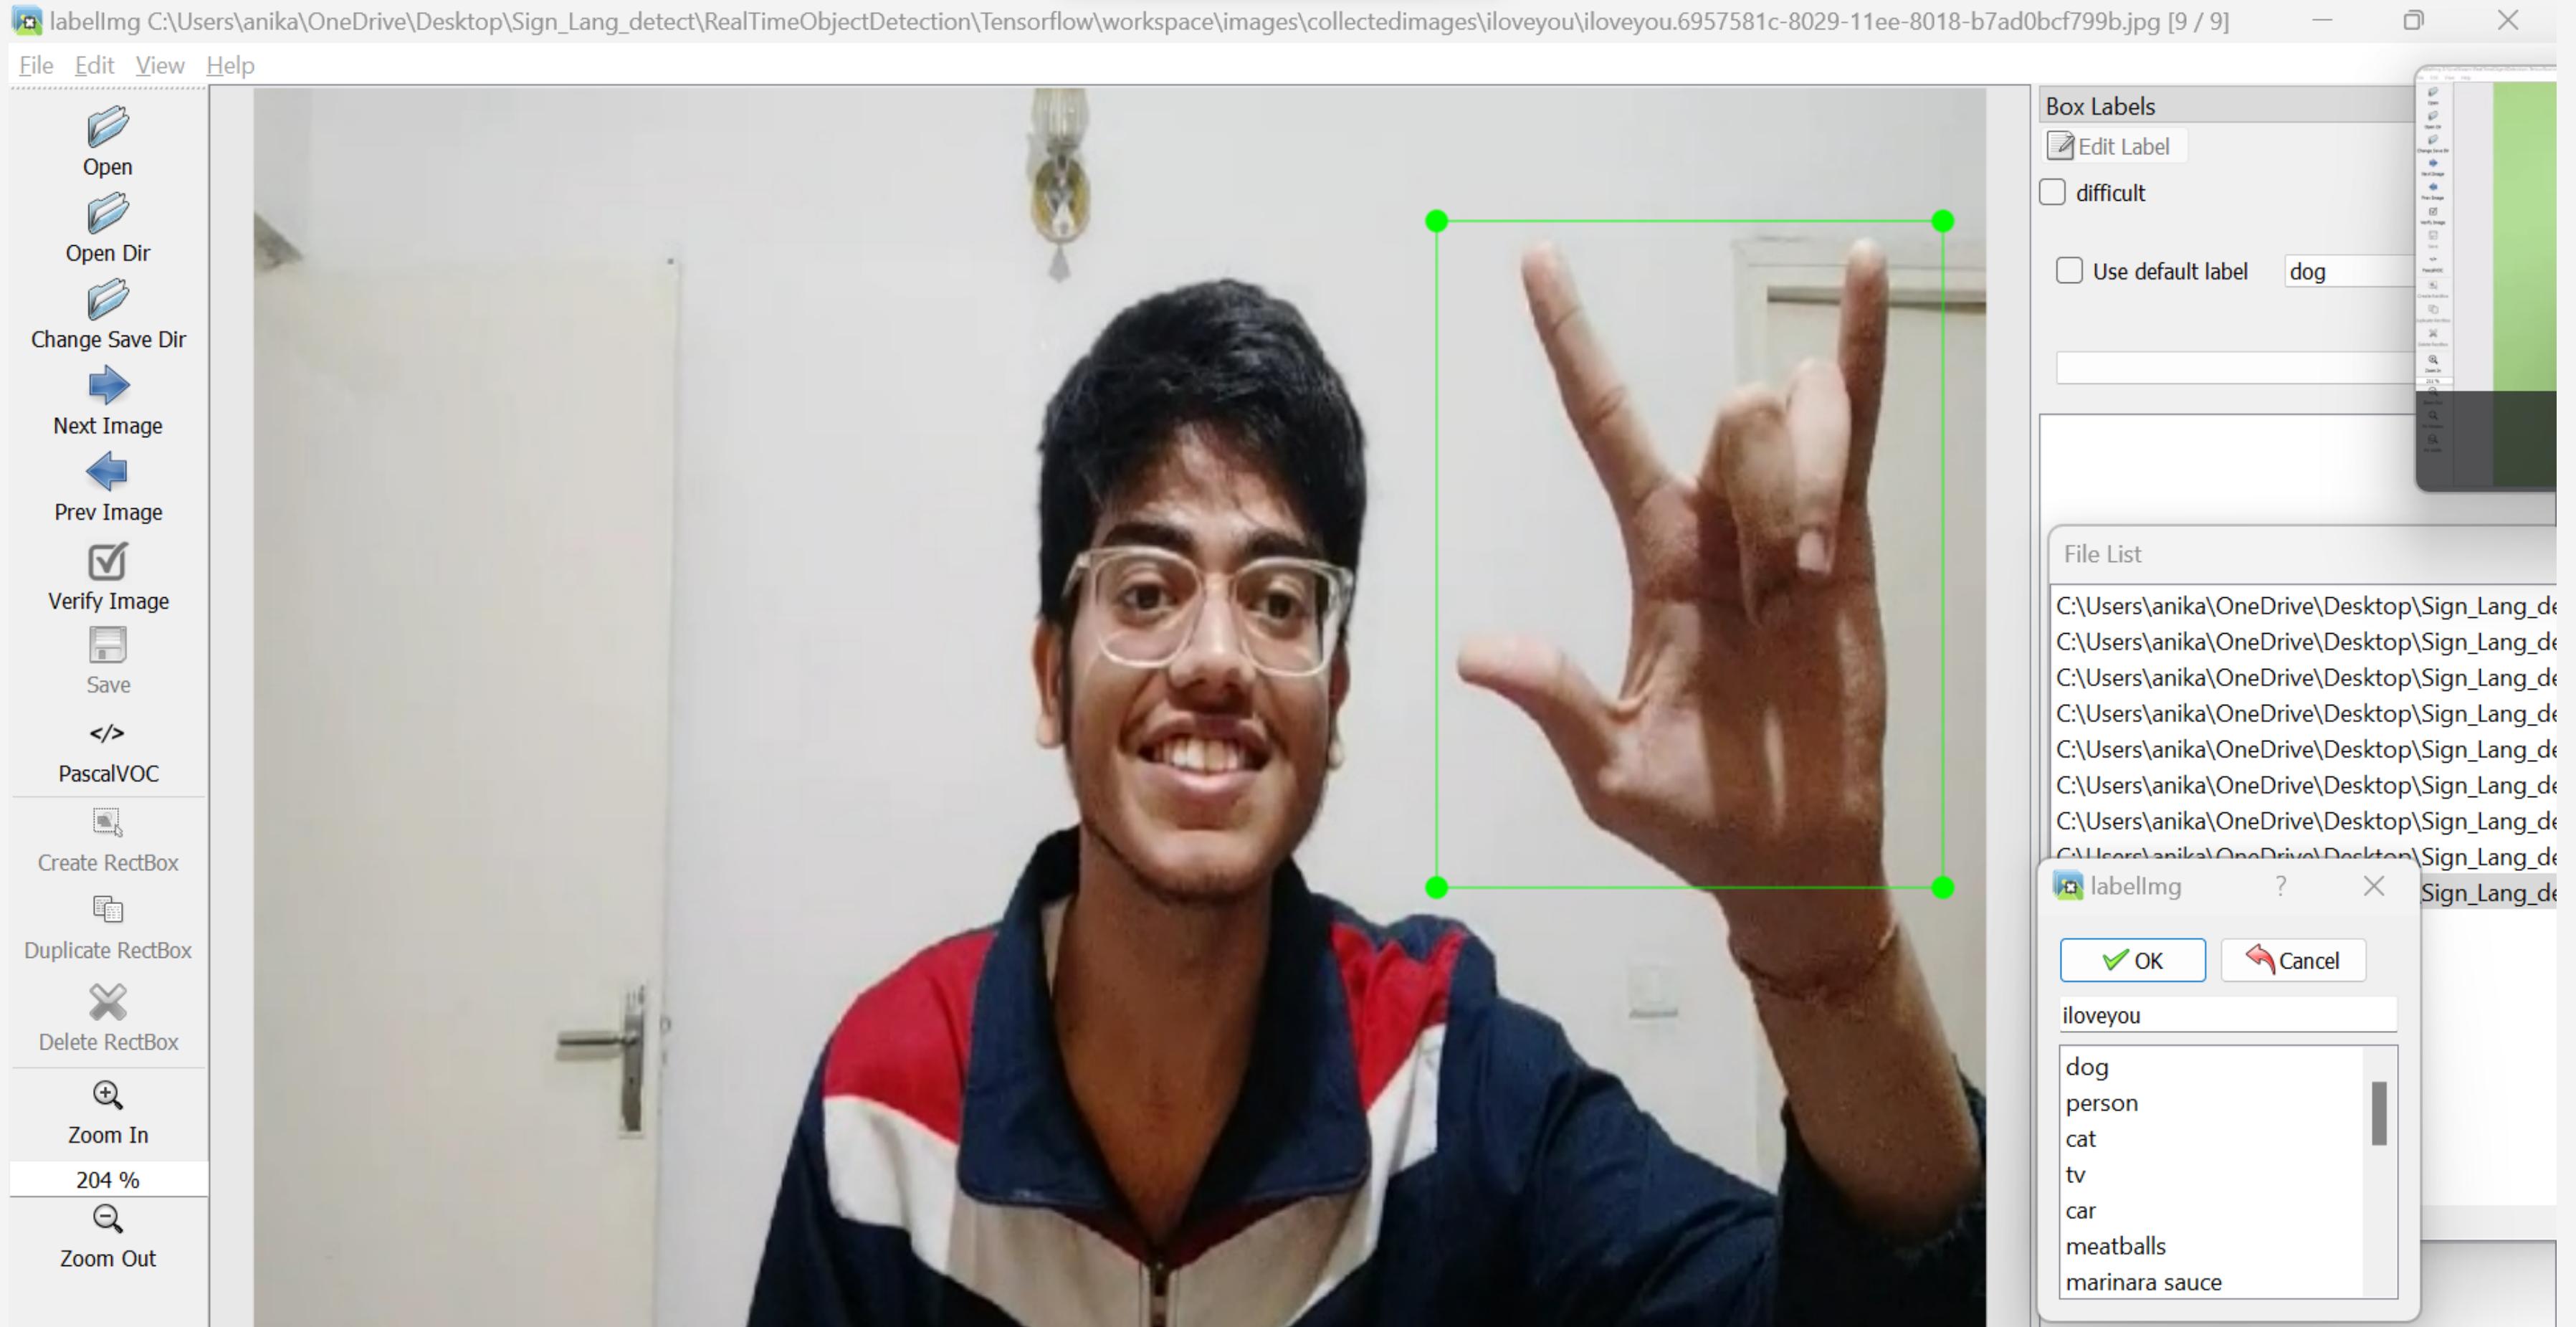
No



Thanks



Labelling Image for Object Detection





Training the model.

Used TensorFlow_2's **SSD_mobilenet_v2** detection model, pre-trained on the [COCO 2017](#) dataset.

Used this pre-trained model for ‘Transfer learning’.

```
(base) C:\Users\anika\OneDrive\Desktop\Sign_Lang_detect\RealTimeObjectDetection>conda activate hh
```

```
(hh) C:\Users\anika\OneDrive\Desktop\Sign_Lang_detect\RealTimeObjectDetection>python Tensorflow\models\research\object_detection\model_main_tf2.py --model_dir=Tensorflow\workspace\models\my_ssd_mobnet --pipeline_config_path=Tensorflow\workspace\models\my_ssd_mobnet/pipeline.config --num_train_steps=10000  
C:\Users\anika\anaconda3\envs\hh\lib\site-packages\tensorflow_addons\utils\tfa_eol_msg.py:23: UserWarning:
```

TensorFlow Addons (TFA) has ended development and introduction of new features.

TFA has entered a minimal maintenance and release mode until a planned end of life in May 2024.

Please modify downstream libraries to take dependencies from other repositories in our TensorFlow community (e.g. Keras, Keras-CV, and Keras-NLP).

For more information see: <https://github.com/tensorflow/addons/issues/2807>

```
warnings.warn(
```

```
C:\Users\anika\anaconda3\envs\hh\lib\site-packages\tensorflow_addons\utils\ensure_tf_install.py:53: UserWarning: TensorFlow Addons supports using Python ops for all Tensorflow versions above or equal to 2.12.0 and strictly below 2.15.0 (nightly versions are not supported).
```

The versions of TensorFlow you are currently using is 2.10.1 and is not supported.

Some things might work, some things might not.

If you were to encounter a bug, do not file an issue.

If you want to make sure you're using a tested and supported configuration, either change the TensorFlow version or the TensorFlow Addons's version.

You can find the compatibility matrix in TensorFlow Addon's readme:

<https://github.com/tensorflow/addons>

```
warnings.warn(
```

```
2023-11-15 11:04:42.825036: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX AVX2
```

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

Anaconda Prompt (anaconda) X + v

You can find the compatibility matrix in TensorFlow Addon's readme:
<https://github.com/tensorflow/addons>

```
warnings.warn(  
2023-11-15 11:04:42.825036: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized  
with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical oper  
ations: AVX AVX2  
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.  
2023-11-15 11:04:44.263178: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1616] Created device /job:localhost/r  
eplica:0/task:0/device:GPU:0 with 1340 MB memory: -> device: 0, name: NVIDIA GeForce MX250, pci bus id: 0000:01:00.  
0, compute capability: 6.1  
INFO:tensorflow:Using MirroredStrategy with devices ('/job:localhost/replica:0/task:0/device:GPU:0',)  
I1115 11:04:44.409381 15804 mirrored_strategy.py:374] Using MirroredStrategy with devices ('/job:localhost/replica:0  
/task:0/device:GPU:0',)  
INFO:tensorflow:Maybe overwriting train_steps: 10000  
I1115 11:04:44.446255 15804 config_util.py:552] Maybe overwriting train_steps: 10000  
INFO:tensorflow:Maybe overwriting use_bfloat16: False  
I1115 11:04:44.446255 15804 config_util.py:552] Maybe overwriting use_bfloat16: False  
WARNING:tensorflow:From C:\Users\anika\anaconda3\envs\hh\lib\site-packages\object_detection\model_lib_v2.py:563: Str  
ategyBase.experimental_distribute_datasets_from_function (from tensorflow.python.distribute.distribute_lib) is depre  
cated and will be removed in a future version.  
Instructions for updating:  
rename to distribute_datasets_from_function  
W1115 11:04:44.483953 15804 deprecation.py:350] From C:\Users\anika\anaconda3\envs\hh\lib\site-packages\object_detec  
tion\model_lib_v2.py:563: StrategyBase.experimental_distribute_datasets_from_function (from tensorflow.python.distri  
bute.distribute_lib) is deprecated and will be removed in a future version.  
Instructions for updating:  
rename to distribute_datasets_from_function  
INFO:tensorflow:Reading unweighted datasets: ['Tensorflow/workspace/annotations/train.record']  
I1115 11:04:44.499662 15804 dataset_builder.py:162] Reading unweighted datasets: ['Tensorflow/workspace/annotations/  
train.record']
```

memory trying to allocate 983.51MiB with freed_by_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

2023-11-15 11:06:39.619296: W tensorflow/core/common_runtime/bfc_allocator.cc:290] Allocator (GPU_0_bfc) ran out of memory trying to allocate 827.50MiB with freed_by_count=0. The caller indicates that this is not a failure, but this may mean that there could be performance gains if more memory were available.

INFO:tensorflow:Step 100 per-step time 0.845s

I1115 11:07:09.494543 15804 model_lib_v2.py:705] Step 100 per-step time 0.845s

INFO:tensorflow:{'Loss/classification_loss': 0.3461144,
'Loss/localization_loss': 0.32647595,
'Loss/regularization_loss': 0.15391405,
'Loss/total_loss': 0.8265044,
'learning_rate': 0.0319994}

I1115 11:07:09.529289 15804 model_lib_v2.py:708] {'Loss/classification_loss': 0.3461144,
'Loss/localization_loss': 0.32647595,
'Loss/regularization_loss': 0.15391405,
'Loss/total_loss': 0.8265044,
'learning_rate': 0.0319994}

INFO:tensorflow:Step 200 per-step time 0.386s

I1115 11:07:48.121766 15804 model_lib_v2.py:705] Step 200 per-step time 0.386s

INFO:tensorflow:{'Loss/classification_loss': 0.2539425,
'Loss/localization_loss': 0.20083773,
'Loss/regularization_loss': 0.15383898,
'Loss/total_loss': 0.6086192,
'learning_rate': 0.0373328}

I1115 11:07:48.123761 15804 model_lib_v2.py:708] {'Loss/classification_loss': 0.2539425,
'Loss/localization_loss': 0.20083773,
'Loss/regularization_loss': 0.15383898,
'Loss/total_loss': 0.6086192,
'learning_rate': 0.0373328}

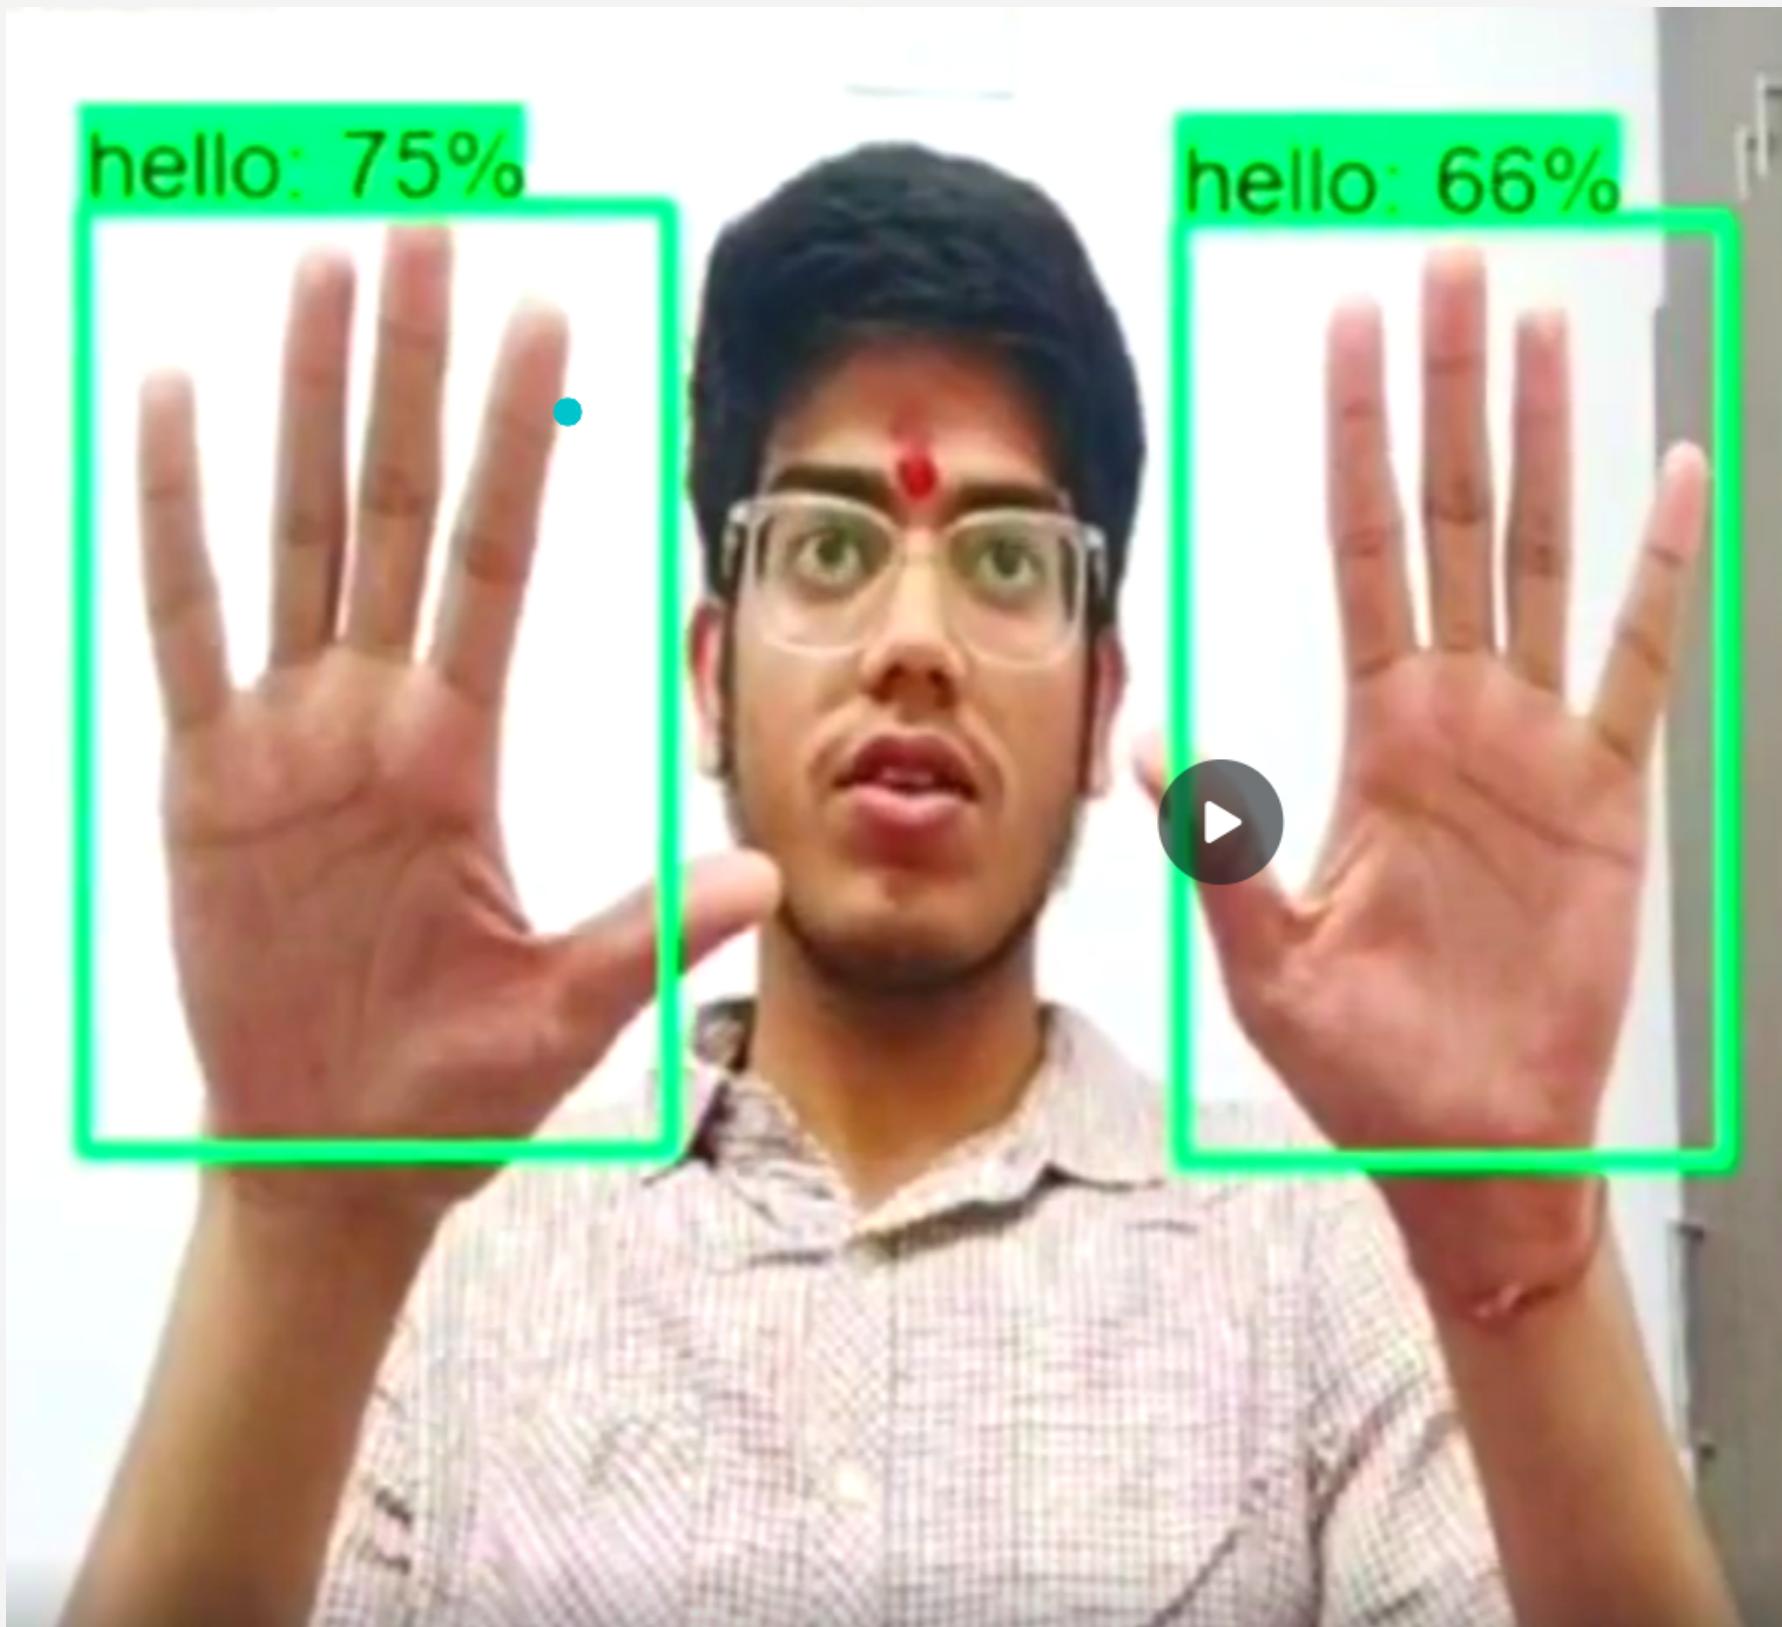
INFO:tensorflow:Step 300 per-step time 0.407s

Anaconda Prompt (anaconda) X + ▾

```
'Loss/localization_loss': 0.010172847,
'Loss/regularization_loss': 0.09720672,
'Loss/total_loss': 0.13367103,
'learning_rate': 0.07380057}
INFO:tensorflow:Step 9900 per-step time 0.514s
I1115 12:29:41.151496 15804 model_lib_v2.py:705] Step 9900 per-step time 0.514s
INFO:tensorflow:{'Loss/classification_loss': 0.05131547,
'Loss/localization_loss': 0.009596345,
'Loss/regularization_loss': 0.09673182,
'Loss/total_loss': 0.15764363,
'learning_rate': 0.073662736}
I1115 12:29:41.154494 15804 model_lib_v2.py:708] {'Loss/classification_loss': 0.05131547,
'Loss/localization_loss': 0.009596345,
'Loss/regularization_loss': 0.09673182,
'Loss/total_loss': 0.15764363,
'learning_rate': 0.073662736}
INFO:tensorflow:Step 10000 per-step time 0.488s
I1115 12:30:29.914326 15804 model_lib_v2.py:705] Step 10000 per-step time 0.488s
INFO:tensorflow:{'Loss/classification_loss': 0.07445922,
'Loss/localization_loss': 0.018851975,
'Loss/regularization_loss': 0.09630479,
'Loss/total_loss': 0.18961598,
'learning_rate': 0.07352352}
I1115 12:30:29.917315 15804 model_lib_v2.py:708] {'Loss/classification_loss': 0.07445922,
'Loss/localization_loss': 0.018851975,
'Loss/regularization_loss': 0.09630479,
'Loss/total_loss': 0.18961598,
'learning_rate': 0.07352352}
```



Results



Hello



YES



I Love You



No



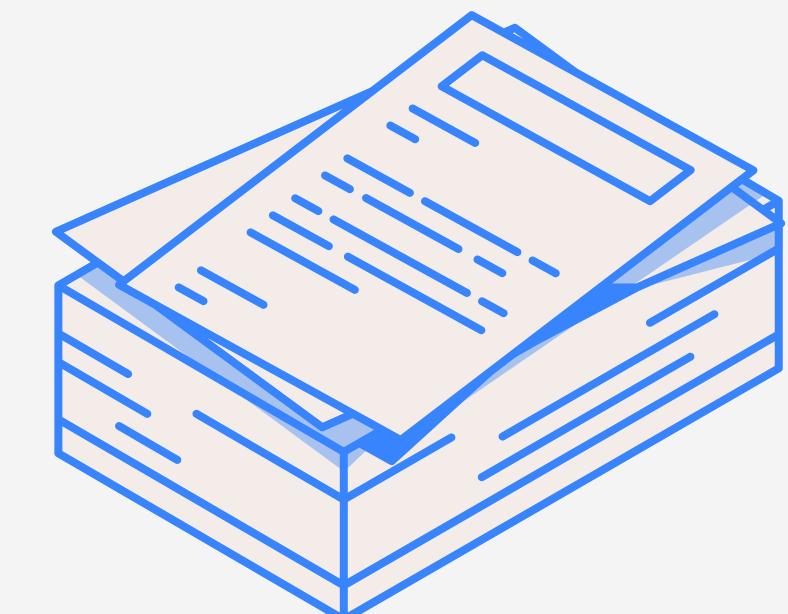
Thanks

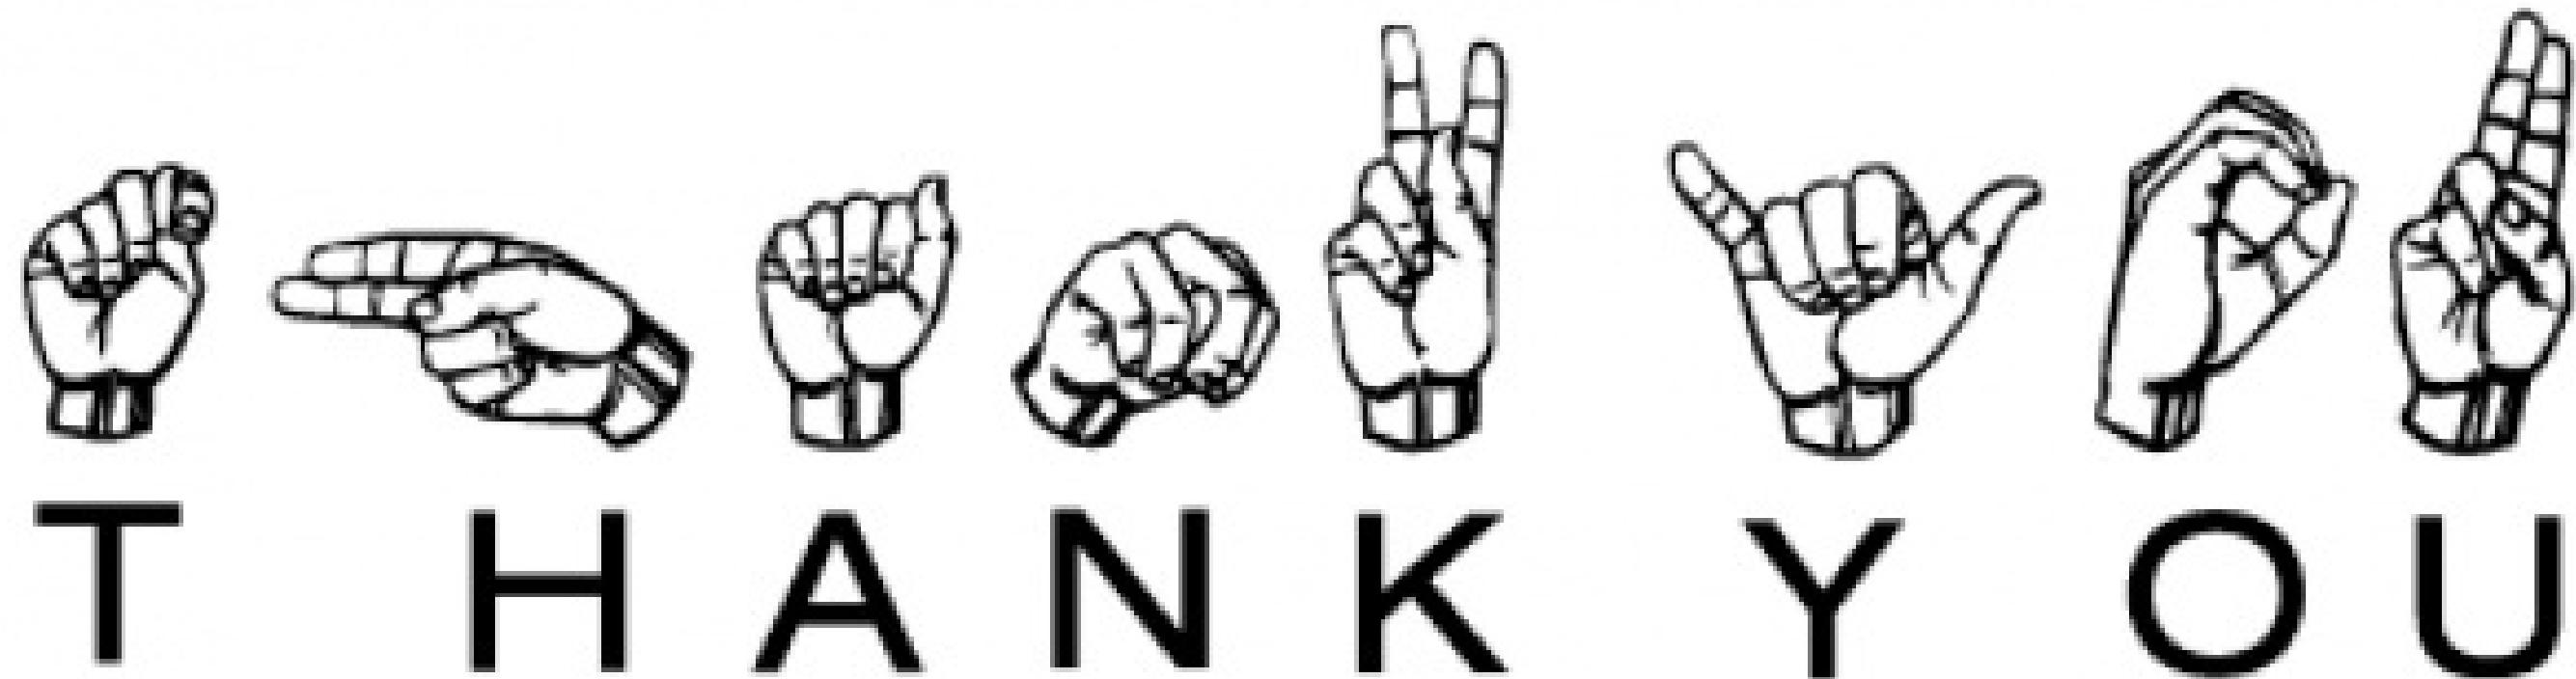
Future Plans.

We plan to combine the two approaches demonstarted in the presentation to develop an end-to-end deployed AI system which can perform Sign Language Detection in real-time.

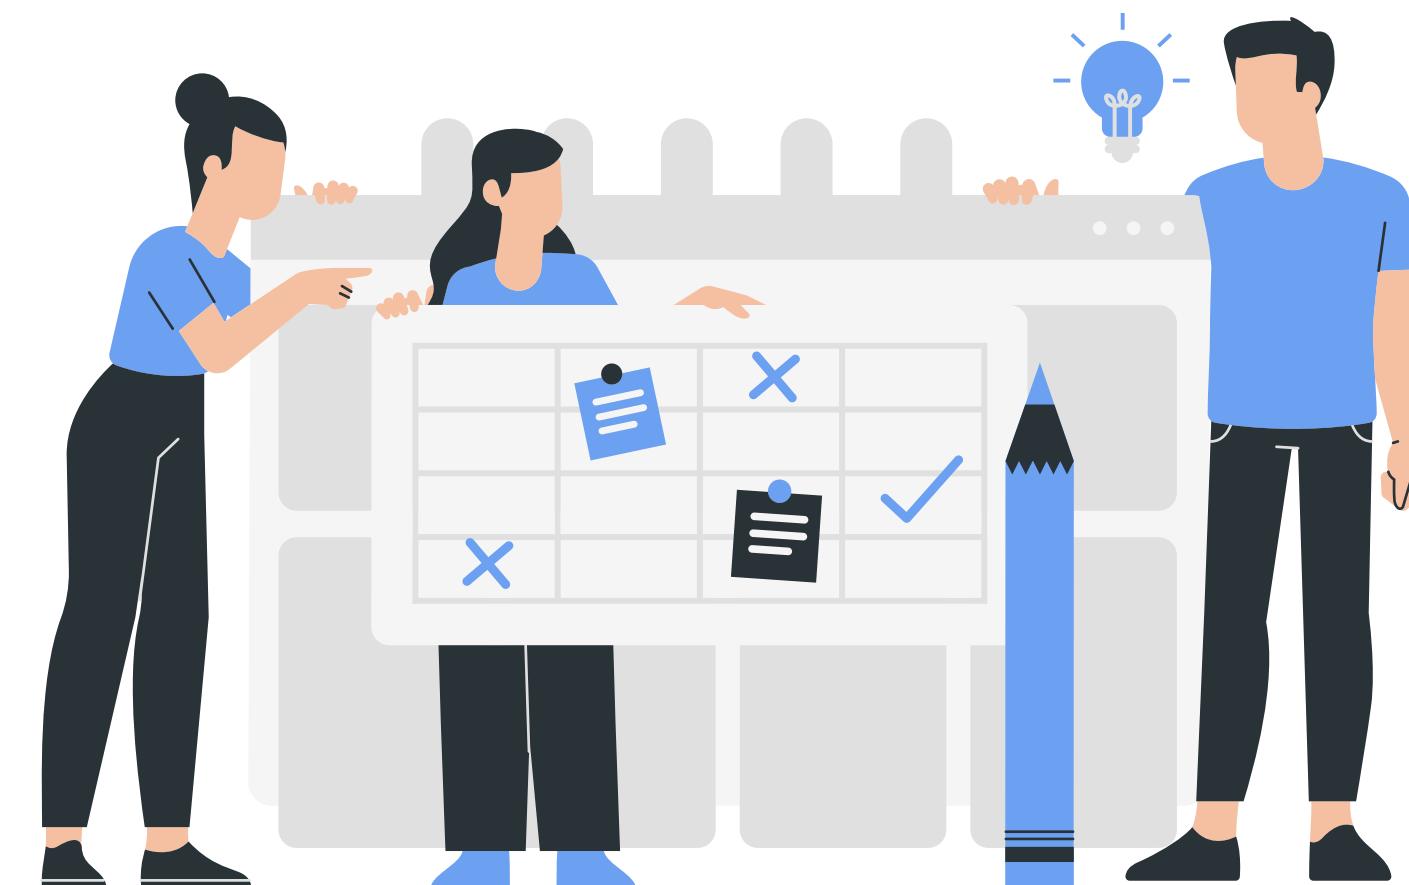
References

- American sign language (ASL) recognition based on Hough Transform and Neural Networks. Qutaishat Munib, 2007.
- MacMaster, Gordon, "Sign Language Translation Using Machine Learning and Computer Vision" (2020).UVM Honors College Senior Theses. 480.
- https://www.youtube.com/watch?v=V0Pk_dPU2IY&t=4228s
- [https://youtube.com/playlist=PLZoTAELRMXVNvTfHyJxPRcQkpV8ubBwH o&si=Nrrug5F3zotxGSkZ](https://youtube.com/playlist=PLZoTAELRMXVNvTfHyJxPRcQkpV8ubBwHo&si=Nrrug5F3zotxGSkZ)
- <https://www.kaggle.com/datasets/debashishsau/aslamerican-sign-language-aplhabet-dataset>





Yashi Sharma (58)



Purnima Kumar (41)

Harsh Yadav (20)